

 Search the docs ...

[Installation](#)

[Package overview](#)

[Getting started tutorials](#)

[Comparison with other tools](#)

[Comparison with R / R libraries](#)

[Comparison with SQL](#)

[Comparison with SAS](#)

[Comparison with Stata](#)

[Community tutorials](#)

Comparison with SQL

Since many potential pandas users have some familiarity with [SQL](#), this page is meant to provide some examples of how various SQL operations would be performed using pandas.

If you're new to pandas, you might want to first read through [10 Minutes to pandas](#) to familiarize yourself with the library.

As is customary, we import pandas and NumPy as follows:

```
In [1]: import pandas as pd
In [2]: import numpy as np
```

Most of the examples will utilize the `tips` dataset found within pandas tests. We'll read the data into a DataFrame called `tips` and assume we have a database table of the same name and structure.

```
In [3]: url = ('https://raw.githubusercontent.com/pandas-dev/
...:         '/pandas/master/pandas/tests/io/data/csv/tips.csv')
...:

In [4]: tips = pd.read_csv(url)

In [5]: tips.head()
Out[5]:
   total_bill  tip  sex smoker  day  time  size
0      16.99  1.01 Female    No  Sun  Dinner     2
1      10.34  1.66   Male    No  Sun  Dinner     3
2      21.01  3.50   Male    No  Sun  Dinner     3
3      23.68  3.31   Male    No  Sun  Dinner     2
4      24.59  3.61 Female    No  Sun  Dinner     4
```

SELECT

In SQL, selection is done using a comma-separated list of columns you'd like to select (or a `*` to select all columns):

```
SELECT total_bill, tip, smoker, time
FROM tips
LIMIT 5;
```

With pandas, column selection is done by passing a list of column names to your DataFrame:

```
In [6]: tips[['total_bill', 'tip', 'smoker', 'time']].head(5)
Out[6]:
   total_bill  tip smoker  time
0      16.99  1.01    No  Dinner
1      10.34  1.66    No  Dinner
2      21.01  3.50    No  Dinner
3      23.68  3.31    No  Dinner
4      24.59  3.61    No  Dinner
```

Calling the DataFrame without the list of column names would display all columns (akin to SQL's `*`).

In SQL, you can add a calculated column:

```
SELECT *, tip/total_bill as tip_rate
FROM tips
LIMIT 5;
```

With pandas, you can use the [DataFrame.assign\(\)](#) method of a DataFrame to append a new column:

```
In [7]: tips.assign(tip_rate=tips['tip'] / tips['total_bill']).head(5)
Out[7]:
```

| | total_bill | tip | sex | smoker | day | time | size | tip_rate |
|---|------------|------|--------|--------|-----|--------|------|----------|
| 0 | 16.99 | 1.01 | Female | No | Sun | Dinner | 2 | 0.059447 |
| 1 | 10.34 | 1.66 | Male | No | Sun | Dinner | 3 | 0.160542 |
| 2 | 21.01 | 3.50 | Male | No | Sun | Dinner | 3 | 0.166587 |
| 3 | 23.68 | 3.31 | Male | No | Sun | Dinner | 2 | 0.139780 |
| 4 | 24.59 | 3.61 | Female | No | Sun | Dinner | 4 | 0.146808 |

WHERE

Filtering in SQL is done via a WHERE clause.

```
SELECT *
FROM tips
WHERE time = 'Dinner'
LIMIT 5;
```

DataFrames can be filtered in multiple ways; the most intuitive of which is using [boolean indexing](#)

```
In [8]: tips[tips['time'] == 'Dinner'].head(5)
Out[8]:
```

| | total_bill | tip | sex | smoker | day | time | size |
|---|------------|------|--------|--------|-----|--------|------|
| 0 | 16.99 | 1.01 | Female | No | Sun | Dinner | 2 |
| 1 | 10.34 | 1.66 | Male | No | Sun | Dinner | 3 |
| 2 | 21.01 | 3.50 | Male | No | Sun | Dinner | 3 |
| 3 | 23.68 | 3.31 | Male | No | Sun | Dinner | 2 |
| 4 | 24.59 | 3.61 | Female | No | Sun | Dinner | 4 |

The above statement is simply passing a **Series** of True/False objects to the DataFrame, returning all rows with True.

```
In [9]: is_dinner = tips['time'] == 'Dinner'

In [10]: is_dinner.value_counts()
Out[10]:
True      176
False      68
Name: time, dtype: int64

In [11]: tips[is_dinner].head(5)
Out[11]:
```

| | total_bill | tip | sex | smoker | day | time | size |
|---|------------|------|--------|--------|-----|--------|------|
| 0 | 16.99 | 1.01 | Female | No | Sun | Dinner | 2 |
| 1 | 10.34 | 1.66 | Male | No | Sun | Dinner | 3 |
| 2 | 21.01 | 3.50 | Male | No | Sun | Dinner | 3 |
| 3 | 23.68 | 3.31 | Male | No | Sun | Dinner | 2 |
| 4 | 24.59 | 3.61 | Female | No | Sun | Dinner | 4 |

Just like SQL’s OR and AND, multiple conditions can be passed to a DataFrame using | (OR) and & (AND).

```
-- tips of more than $5.00 at Dinner meals
SELECT *
FROM tips
WHERE time = 'Dinner' AND tip > 5.00;
```

```
# tips of more than $5.00 at Dinner meals
In [12]: tips[(tips['time'] == 'Dinner') & (tips['tip'] > 5.00)]
Out[12]:
```

| | total_bill | tip | sex | smoker | day | time | size |
|-----|------------|-------|--------|--------|-----|--------|------|
| 23 | 39.42 | 7.58 | Male | No | Sat | Dinner | 4 |
| 44 | 30.40 | 5.60 | Male | No | Sun | Dinner | 4 |
| 47 | 32.40 | 6.00 | Male | No | Sun | Dinner | 4 |
| 52 | 34.81 | 5.20 | Female | No | Sun | Dinner | 4 |
| 59 | 48.27 | 6.73 | Male | No | Sat | Dinner | 4 |
| 116 | 29.93 | 5.07 | Male | No | Sun | Dinner | 4 |
| 155 | 29.85 | 5.14 | Female | No | Sun | Dinner | 5 |
| 170 | 50.81 | 10.00 | Male | Yes | Sat | Dinner | 3 |
| 172 | 7.25 | 5.15 | Male | Yes | Sun | Dinner | 2 |
| 181 | 23.33 | 5.65 | Male | Yes | Sun | Dinner | 2 |
| 183 | 23.17 | 6.50 | Male | Yes | Sun | Dinner | 4 |
| 211 | 25.89 | 5.16 | Male | Yes | Sat | Dinner | 4 |
| 212 | 48.33 | 9.00 | Male | No | Sat | Dinner | 4 |
| 214 | 28.17 | 6.50 | Female | Yes | Sat | Dinner | 3 |
| 239 | 29.03 | 5.92 | Male | No | Sat | Dinner | 3 |

```
-- tips by parties of at least 5 diners OR bill total was more than $45
SELECT *
FROM tips
WHERE size >= 5 OR total_bill > 45;
```

```
# tips by parties of at least 5 diners OR bill total was more than $45
In [13]: tips[(tips['size'] >= 5) | (tips['total_bill'] > 45)]
Out[13]:
```

| | total_bill | tip | sex | smoker | day | time | size |
|-----|------------|-------|--------|--------|------|--------|------|
| 59 | 48.27 | 6.73 | Male | No | Sat | Dinner | 4 |
| 125 | 29.80 | 4.20 | Female | No | Thur | Lunch | 6 |
| 141 | 34.30 | 6.70 | Male | No | Thur | Lunch | 6 |
| 142 | 41.19 | 5.00 | Male | No | Thur | Lunch | 5 |
| 143 | 27.05 | 5.00 | Female | No | Thur | Lunch | 6 |
| 155 | 29.85 | 5.14 | Female | No | Sun | Dinner | 5 |
| 156 | 48.17 | 5.00 | Male | No | Sun | Dinner | 6 |
| 170 | 50.81 | 10.00 | Male | Yes | Sat | Dinner | 3 |
| 182 | 45.35 | 3.50 | Male | Yes | Sun | Dinner | 3 |
| 185 | 20.69 | 5.00 | Male | No | Sun | Dinner | 5 |
| 187 | 30.46 | 2.00 | Male | Yes | Sun | Dinner | 5 |
| 212 | 48.33 | 9.00 | Male | No | Sat | Dinner | 4 |
| 216 | 28.15 | 3.00 | Male | Yes | Sat | Dinner | 5 |

NULL checking is done using the [notna\(\)](#) and [isna\(\)](#) methods.

```
In [14]: frame = pd.DataFrame({'col1': ['A', 'B', np.NaN, 'C', 'D'],
.....:                        'col2': ['F', np.NaN, 'G', 'H', 'I']})
.....:

In [15]: frame
Out[15]:
```

| | col1 | col2 |
|---|------|------|
| 0 | A | F |
| 1 | B | NaN |
| 2 | NaN | G |
| 3 | C | H |
| 4 | D | I |

Assume we have a table of the same structure as our DataFrame above. We can see only the records where `col2` IS NULL with the following query:

```
SELECT *
FROM frame
WHERE col2 IS NULL;
```

```
In [16]: frame[frame['col2'].isna()]
Out[16]:
```

| | col1 | col2 |
|---|------|------|
| 1 | B | NaN |

Getting items where `col1` IS NOT NULL can be done with [notna\(\)](#).

```
SELECT *
FROM frame
WHERE col1 IS NOT NULL;
```

```
In [17]: frame[frame['col1'].notna()]
Out[17]:
```

| | col1 | col2 |
|---|------|------|
| 0 | A | F |
| 1 | B | NaN |
| 3 | C | H |
| 4 | D | I |

GROUP BY

In pandas, SQL's GROUP BY operations are performed using the similarly named [groupby\(\)](#) method. [groupby\(\)](#) typically refers to a process where we'd like to split a dataset into groups, apply some function (typically aggregation) , and then combine the groups together.

A common SQL operation would be getting the count of records in each group throughout a dataset. For instance, a query getting us the number of tips left by sex:

```
SELECT sex, count(*)
FROM tips
GROUP BY sex;
/*
Female      87
Male       157
*/
```

The pandas equivalent would be:

```
In [18]: tips.groupby('sex').size()
Out[18]:
sex
Female      87
Male       157
dtype: int64
```

Notice that in the pandas code we used `size()` and not `count()`. This is because `count()` applies the function to each column, returning the number of `not null` records within each.

```
In [19]: tips.groupby('sex').count()
Out[19]:
      total_bill  tip  smoker  day  time  size
sex
Female         87   87      87   87    87    87
Male         157  157     157  157   157   157
```

Alternatively, we could have applied the `count()` method to an individual column:

```
In [20]: tips.groupby('sex')['total_bill'].count()
Out[20]:
sex
Female      87
Male       157
Name: total_bill, dtype: int64
```

Multiple functions can also be applied at once. For instance, say we'd like to see how tip amount differs by day of the week - `agg()` allows you to pass a dictionary to your grouped DataFrame, indicating which functions to apply to specific columns.

```
SELECT day, AVG(tip), COUNT(*)
FROM tips
GROUP BY day;
/*
Fri  2.734737  19
Sat  2.993103  87
Sun  3.255132  76
Thur 2.771452  62
*/
```

```
In [21]: tips.groupby('day').agg({'tip': np.mean, 'day': np.size})
Out[21]:
      tip  day
day
Fri  2.734737  19
Sat  2.993103  87
Sun  3.255132  76
Thur 2.771452  62
```

Grouping by more than one column is done by passing a list of columns to the `groupby()` method.

```
SELECT smoker, day, COUNT(*), AVG(tip)
FROM tips
GROUP BY smoker, day;
/*
smoker day
No      Fri    4  2.812500
        Sat   45  3.102889
        Sun   57  3.167895
        Thur  45  2.673778
Yes     Fri   15  2.714000
        Sat   42  2.875476
        Sun   19  3.516842
        Thur  17  3.030000
*/
```

```
In [22]: tips.groupby(['smoker', 'day']).agg({'tip': [np.size, np.mean]})
Out[22]:
      tip
smoker day size      mean
No      Fri   4.0  2.812500
        Sat  45.0  3.102889
        Sun  57.0  3.167895
        Thur 45.0  2.673778
Yes     Fri  15.0  2.714000
        Sat  42.0  2.875476
        Sun  19.0  3.516842
        Thur 17.0  3.030000
```

JOIN

JOINS can be performed with `join()` or `merge()`. By default, `join()` will join the DataFrames on their indices. Each method has parameters allowing you to specify the type of join to perform (LEFT, RIGHT, INNER, FULL) or the columns to join on (column names or indices).

```
In [23]: df1 = pd.DataFrame({'key': ['A', 'B', 'C', 'D'],
.....:                      'value': np.random.randn(4)})
.....:

In [24]: df2 = pd.DataFrame({'key': ['B', 'D', 'D', 'E'],
.....:                      'value': np.random.randn(4)})
.....:
```

Assume we have two database tables of the same name and structure as our DataFrames.

Now let’s go over the various types of JOINS.

INNER JOIN

```
SELECT *
FROM df1
INNER JOIN df2
  ON df1.key = df2.key;
```

```
# merge performs an INNER JOIN by default
In [25]: pd.merge(df1, df2, on='key')
Out[25]:
   key  value_x  value_y
0    B -0.282863  1.212112
1    D -1.135632 -0.173215
2    D -1.135632  0.119209
```

`merge()` also offers parameters for cases when you’d like to join one DataFrame’s column with another DataFrame’s index.

```
In [26]: indexed_df2 = df2.set_index('key')

In [27]: pd.merge(df1, indexed_df2, left_on='key', right_index=True)
Out[27]:
   key  value_x  value_y
1    B -0.282863  1.212112
3    D -1.135632 -0.173215
3    D -1.135632  0.119209
```

LEFT OUTER JOIN

```
-- show all records from df1
SELECT *
FROM df1
LEFT OUTER JOIN df2
  ON df1.key = df2.key;
```

```
# show all records from df1
In [28]: pd.merge(df1, df2, on='key', how='left')
Out[28]:
   key  value_x  value_y
0    A  0.469112      NaN
1    B -0.282863  1.212112
2    C -1.509059      NaN
3    D -1.135632 -0.173215
4    D -1.135632  0.119209
```

RIGHT JOIN

```
-- show all records from df2
SELECT *
FROM df1
RIGHT OUTER JOIN df2
  ON df1.key = df2.key;
```

```
# show all records from df2
In [29]: pd.merge(df1, df2, on='key', how='right')
Out[29]:
```

| | key | value_x | value_y |
|---|-----|-----------|-----------|
| 0 | B | -0.282863 | 1.212112 |
| 1 | D | -1.135632 | -0.173215 |
| 2 | D | -1.135632 | 0.119209 |
| 3 | E | NaN | -1.044236 |

FULL JOIN

pandas also allows for FULL JOINS, which display both sides of the dataset, whether or not the joined columns find a match. As of writing, FULL JOINS are not supported in all RDBMS (MySQL).

```
-- show all records from both tables
SELECT *
FROM df1
FULL OUTER JOIN df2
ON df1.key = df2.key;
```

```
# show all records from both frames
In [30]: pd.merge(df1, df2, on='key', how='outer')
Out[30]:
```

| | key | value_x | value_y |
|---|-----|-----------|-----------|
| 0 | A | 0.469112 | NaN |
| 1 | B | -0.282863 | 1.212112 |
| 2 | C | -1.509059 | NaN |
| 3 | D | -1.135632 | -0.173215 |
| 4 | D | -1.135632 | 0.119209 |
| 5 | E | NaN | -1.044236 |

UNION

UNION ALL can be performed using [concat\(\)](#).

```
In [31]: df1 = pd.DataFrame({'city': ['Chicago', 'San Francisco', 'New York City'],
.....:                      'rank': range(1, 4)})
.....:

In [32]: df2 = pd.DataFrame({'city': ['Chicago', 'Boston', 'Los Angeles'],
.....:                      'rank': [1, 4, 5]})
.....:
```

```
SELECT city, rank
FROM df1
UNION ALL
SELECT city, rank
FROM df2;
/*
      city  rank
Chicago    1
San Francisco  2
New York City  3
Chicago    1
Boston      4
Los Angeles  5
*/
```

```
In [33]: pd.concat([df1, df2])
Out[33]:
```

| | city | rank |
|---|---------------|------|
| 0 | Chicago | 1 |
| 1 | San Francisco | 2 |
| 2 | New York City | 3 |
| 0 | Chicago | 1 |
| 1 | Boston | 4 |
| 2 | Los Angeles | 5 |

SQL's UNION is similar to UNION ALL, however UNION will remove duplicate rows.

```
SELECT city, rank
FROM df1
UNION
SELECT city, rank
FROM df2;
-- notice that there is only one Chicago record this time
/*
      city  rank
    Chicago    1
San Francisco    2
New York City    3
      Boston    4
    Los Angeles    5
*/
```

In pandas, you can use `concat()` in conjunction with `drop_duplicates()`.

```
In [34]: pd.concat([df1, df2]).drop_duplicates()
Out[34]:
```

| | city | rank |
|---|---------------|------|
| 0 | Chicago | 1 |
| 1 | San Francisco | 2 |
| 2 | New York City | 3 |
| 1 | Boston | 4 |
| 2 | Los Angeles | 5 |

pandas equivalents for some SQL analytic and aggregate functions

Top n rows with offset

```
-- MySQL
SELECT * FROM tips
ORDER BY tip DESC
LIMIT 10 OFFSET 5;
```

```
In [35]: tips.nlargest(10 + 5, columns='tip').tail(10)
Out[35]:
```

| | total_bill | tip | sex | smoker | day | time | size |
|-----|------------|------|--------|--------|------|--------|------|
| 183 | 23.17 | 6.50 | Male | Yes | Sun | Dinner | 4 |
| 214 | 28.17 | 6.50 | Female | Yes | Sat | Dinner | 3 |
| 47 | 32.40 | 6.00 | Male | No | Sun | Dinner | 4 |
| 239 | 29.03 | 5.92 | Male | No | Sat | Dinner | 3 |
| 88 | 24.71 | 5.85 | Male | No | Thur | Lunch | 2 |
| 181 | 23.33 | 5.65 | Male | Yes | Sun | Dinner | 2 |
| 44 | 30.40 | 5.60 | Male | No | Sun | Dinner | 4 |
| 52 | 34.81 | 5.20 | Female | No | Sun | Dinner | 4 |
| 85 | 34.83 | 5.17 | Female | No | Thur | Lunch | 4 |
| 211 | 25.89 | 5.16 | Male | Yes | Sat | Dinner | 4 |

Top n rows per group

```
-- Oracle's ROW_NUMBER() analytic function
SELECT * FROM (
  SELECT
    t.*,
    ROW_NUMBER() OVER(PARTITION BY day ORDER BY total_bill DESC) AS rn
  FROM tips t
)
WHERE rn < 3
ORDER BY day, rn;
```

```
In [36]: (tips.assign(rn=tips.sort_values(['total_bill'], ascending=False)
.....:               .groupby(['day'])
.....:               .cumcount() + 1)
.....:       .query('rn < 3')
.....:       .sort_values(['day', 'rn']))
Out[36]:
```

| | total_bill | tip | sex | smoker | day | time | size | rn |
|-----|------------|-------|--------|--------|------|--------|------|----|
| 95 | 40.17 | 4.73 | Male | Yes | Fri | Dinner | 4 | 1 |
| 90 | 28.97 | 3.00 | Male | Yes | Fri | Dinner | 2 | 2 |
| 170 | 50.81 | 10.00 | Male | Yes | Sat | Dinner | 3 | 1 |
| 212 | 48.33 | 9.00 | Male | No | Sat | Dinner | 4 | 2 |
| 156 | 48.17 | 5.00 | Male | No | Sun | Dinner | 6 | 1 |
| 182 | 45.35 | 3.50 | Male | Yes | Sun | Dinner | 3 | 2 |
| 197 | 43.11 | 5.00 | Female | Yes | Thur | Lunch | 4 | 1 |
| 142 | 41.19 | 5.00 | Male | No | Thur | Lunch | 5 | 2 |

the same using `rank(method='first')` function

```
In [37]: (tips.assign(rnk=tips.groupby(['day'])['total_bill']
.....:               .rank(method='first', ascending=False))
.....:       .query('rnk < 3')
.....:       .sort_values(['day', 'rnk']))
Out[37]:
```

| | total_bill | tip | sex | smoker | day | time | size | rnk |
|-----|------------|-------|--------|--------|------|--------|------|-----|
| 95 | 40.17 | 4.73 | Male | Yes | Fri | Dinner | 4 | 1.0 |
| 90 | 28.97 | 3.00 | Male | Yes | Fri | Dinner | 2 | 2.0 |
| 170 | 50.81 | 10.00 | Male | Yes | Sat | Dinner | 3 | 1.0 |
| 212 | 48.33 | 9.00 | Male | No | Sat | Dinner | 4 | 2.0 |
| 156 | 48.17 | 5.00 | Male | No | Sun | Dinner | 6 | 1.0 |
| 182 | 45.35 | 3.50 | Male | Yes | Sun | Dinner | 3 | 2.0 |
| 197 | 43.11 | 5.00 | Female | Yes | Thur | Lunch | 4 | 1.0 |
| 142 | 41.19 | 5.00 | Male | No | Thur | Lunch | 5 | 2.0 |

```
-- Oracle's RANK() analytic function
SELECT * FROM (
  SELECT
    t.*,
    RANK() OVER(PARTITION BY sex ORDER BY tip) AS rnk
  FROM tips t
  WHERE tip < 2
)
WHERE rnk < 3
ORDER BY sex, rnk;
```

Let’s find tips with (rank < 3) per gender group for (tips < 2). Notice that when using `rank(method='min')` function `rnk_min` remains the same for the same `tip` (as Oracle’s RANK() function)

```
In [38]: (tips[tips['tip'] < 2]
.....:       .assign(rnk_min=tips.groupby(['sex'])['tip']
.....:               .rank(method='min'))
.....:       .query('rnk_min < 3')
.....:       .sort_values(['sex', 'rnk_min']))
Out[38]:
```

| | total_bill | tip | sex | smoker | day | time | size | rnk_min |
|-----|------------|------|--------|--------|-----|--------|------|---------|
| 67 | 3.07 | 1.00 | Female | Yes | Sat | Dinner | 1 | 1.0 |
| 92 | 5.75 | 1.00 | Female | Yes | Fri | Dinner | 2 | 1.0 |
| 111 | 7.25 | 1.00 | Female | No | Sat | Dinner | 1 | 1.0 |
| 236 | 12.60 | 1.00 | Male | Yes | Sat | Dinner | 2 | 1.0 |
| 237 | 32.83 | 1.17 | Male | Yes | Sat | Dinner | 2 | 2.0 |

UPDATE

```
UPDATE tips
SET tip = tip*2
WHERE tip < 2;
```

```
In [39]: tips.loc[tips['tip'] < 2, 'tip'] *= 2
```

DELETE

```
DELETE FROM tips
WHERE tip > 9;
```

In pandas we select the rows that should remain, instead of deleting them

```
In [40]: tips = tips.loc[tips['tip'] <= 9]
```

[<< Comparison with R / R libraries](#)

[Comparison with SAS >>](#)